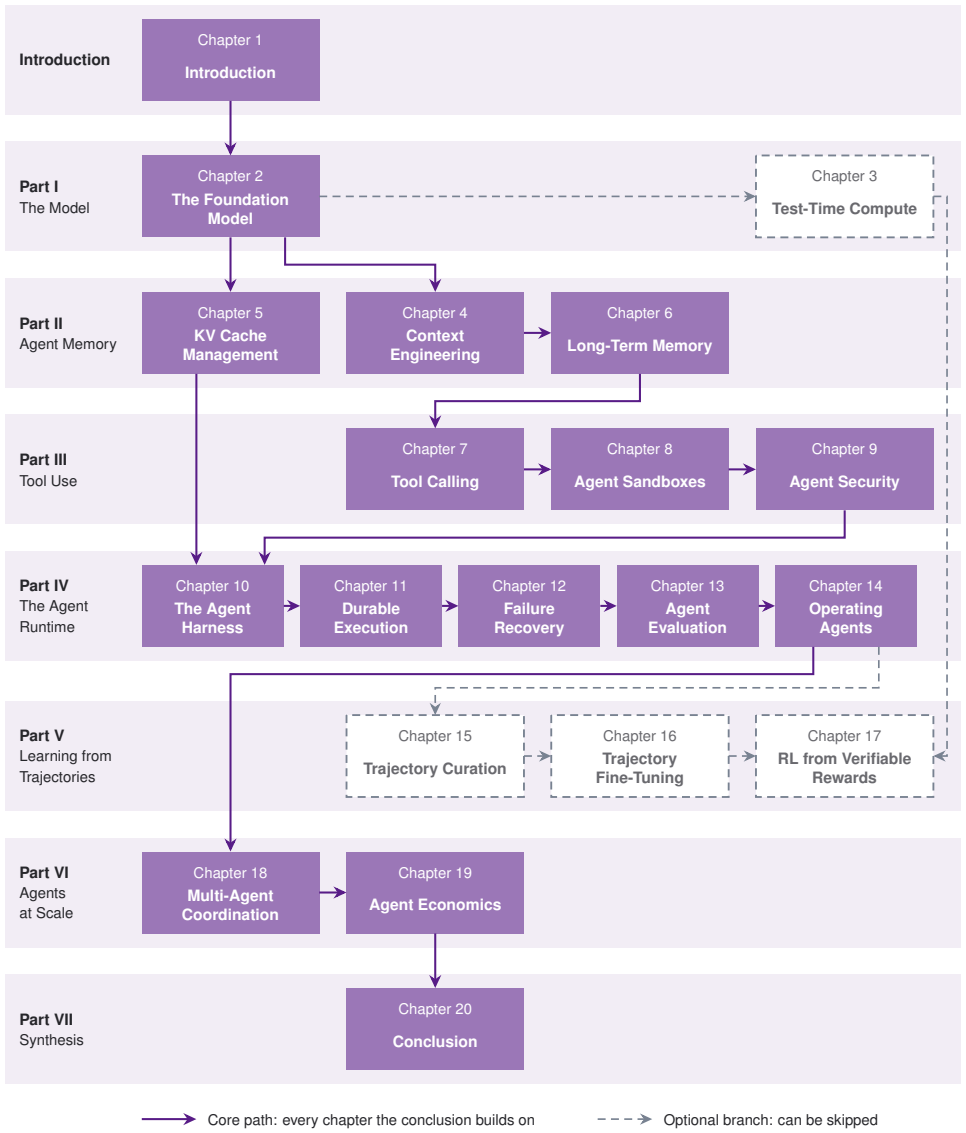
Introduction
Chapter 1
Introduction
Part I
The Model
Chapter 2
The Foundation Model
Chapter 3
Test-Time Compute
Part II
Agent Memory
Chapter 5
KV Cache Management
Chapter 4
Context Engineering
Chapter 6
Long-Term Memory
Part III
Tool Use
Chapter 7
Tool Calling
Chapter 8
Agent Sandboxes
Chapter 9
Agent Security
Part IV
The Agent Runtime
Chapter 10
The Agent Harness
Chapter 11
Durable Execution
Chapter 12
Failure Recovery
Chapter 13
Agent Evaluation
Chapter 14
Operating Agents
Part V
Learning from Trajectories
Chapter 15
Trajectory Curation
Chapter 16
Trajectory Fine-Tuning
Chapter 17
RL from Verifiable Rewards
Part VI
Agents at Scale
Chapter 18
Multi-Agent Coordination
Chapter 19
Agent Economics
Part VII
Synthesis
Chapter 20
Conclusion
Core path: every chapter the conclusion builds on
Optional branch: can be skipped